

Table des matières

	<i>Source.</i>	
	1. Authors	2
	<i>Contributors.</i>	
	2. What is it ?	2
Chapitre 1	Requirements	3
Chapitre 2	Controller Node	4
	1. Preparing Ubuntu	4
	2. Networking	4
	3. MySQL & RabbitMQ	5
	4. RabbitMQ	5
	5. Others	6
	6. Keystone	6
	7. Glance	7
	8. Quantum	8
	9. Nova	9
	10 Cinder	11
	11. Horizon	12
Chapitre 3	Network Node	14
	1. Preparing the Node	14
	2. Networking	15
	3. OpenVSwitch (Part1)	15
	4. Quantum	15
	5. OpenVSwitch (Part2)	17
Chapitre 4	Compute Node	18
	1. Preparing the Node	18
	2. Networking	19
	3. KVM	19
	4. OpenVSwitch	20
	5. Quantum	20
	6. Nova	21
Chapitre 5	Your first VM	24
Chapitre 6	Licensing	26

1. Source

<https://github.com/mseknibilel/OpenStack-Grizzly-Install-Guide>

a Keywords

Multi node, Grizzly, Quantum, Nova, Keystone, Glance, Horizon, Cinder, OpenVSwitch, KVM, Ubuntu Server 12.04/13.04 (64 bits).

1 Authors

Bilel Msekni

1. Contributors

Houssem Medhioub

Djamal Zeghlache

Sandeep Raman

Sam Stoelinga Anil Vishnoi

Gangur Hrushikesh

Wana contribute? Read the guide, send your contribution and get your name listed ;)

2 What is it?

OpenStack Grizzly Install Guide is an easy and tested way to create your own OpenStack platform.

If you like it, don't forget to star it!

Status : Stable

Requirements

Node Role	NICs
Control Node	eth0 (10.10.10.51), eth1 (192.168.100.51)
Network Node	eth0 (10.10.10.52), eth1 (10.20.20.52), eth2 (192.168.100.52)
Compute Node	eth0 (10.10.10.53), eth1 (10.20.20.53)

Note 1 : Always use `dpkg -s` to make sure you are using grizzly packages (version : 2013.1)

Note 2 : This is my current network architecture, you can add as many compute node as you wish.

Controller Node

1 Preparing Ubuntu

- After you install Ubuntu 12.04 or 13.04 Server 64bits, Go in sudo mode and don't leave it until the end of this guide :
`sudo su`
- Add Grizzly repositories [Only for Ubuntu 12.04] :
`apt-get install -y ubuntu-cloud-keyring`
`echo deb http://ubuntu-cloud.archive.canonical.com/ubuntu \`
`precise-updates/grizzly main >> /etc/apt/sources.list.d/grizzly.list`
- Update your system :
`apt-get update -y`
`apt-get upgrade -y`
`apt-get dist-upgrade -y`

2 Networking

- Only one NIC should have an internet access :
`#For Exposing OpenStack API over the internet`
`auto eth1`
`iface eth1 inet static`
`address 192.168.100.51`
`netmask 255.255.255.0`
`gateway 192.168.100.1`
`dns-nameservers 8.8.8.8`
`# Not internet connected(used for OpenStack management)`
`auto eth0`
`iface eth0 inet static`
`address 10.10.10.51`
`netmask 255.255.255.0`
- Restart the networking service :

```
service networking restart
```

3 MySQL & RabbitMQ

- Install MySQL :

```
apt-get install -y mysql-server python-mysqldb
```

- Configure mysql to accept all incoming requests :

```
sed -i 's/127.0.0.1/0.0.0.0/g' /etc/mysql/my.cnf
```

```
service mysql restart
```

- Create these databases :

```
mysql -u root -p
```

```
#Keystone
```

```
CREATE DATABASE keystone;
```

```
GRANT ALL ON keystone.* TO 'keystoneUser'@'%' IDENTIFIED BY 'keystonePass';
```

```
#Glance
```

```
CREATE DATABASE glance;
```

```
GRANT ALL ON glance.* TO 'glanceUser'@'%' IDENTIFIED BY 'glancePass';
```

```
#Quantum
```

```
CREATE DATABASE quantum;
```

```
GRANT ALL ON quantum.* TO 'quantumUser'@'%' IDENTIFIED BY 'quantumPass';
```

```
#Nova
```

```
CREATE DATABASE nova;
```

```
GRANT ALL ON nova.* TO 'novaUser'@'%' IDENTIFIED BY 'novaPass';
```

```
#Cinder
```

```
CREATE DATABASE cinder;
```

```
GRANT ALL ON cinder.* TO 'cinderUser'@'%' IDENTIFIED BY 'cinderPass';
```

```
quit;
```

4 RabbitMQ

- Install RabbitMQ :

```
apt-get install -y rabbitmq-server
```

- Install NTP service :

```
apt-get install -y ntp
```

5 Others

- Install other services :

```
apt-get install -y vlan bridge-utils
```

- Enable IP_Forwarding :

```
sed -i 's/#net.ipv4.ip_forward=1/net.ipv4.ip_forward=1/' \
/etc/sysctl.conf
```

```
# To save you from rebooting, perform the following
sysctl net.ipv4.ip_forward=1
```

6 Keystone

- Start by the keystone packages :

```
apt-get install -y keystone
```

- Adapt the connection attribute in the `/etc/keystone/keystone.conf` to the new database :

```
connection = mysql://keystoneUser:keystonePass@10.10.10.51/keystone
```

- Restart the identity service then synchronize the database :

```
service keystone restart
keystone-manage db_sync
```

- Fill up the keystone database using the two scripts available in the [Scripts folder](#)

Modify the **HOST_IP** and **EXT_HOST_IP** variables before executing the scripts

```
wget https://raw.githubusercontent.com/mseknibilel/OpenStack-Grizzly-Install-Guide/\
OVS_MultiNode/KeystoneScripts/keystone_basic.sh
wget https://raw.githubusercontent.com/mseknibilel/OpenStack-Grizzly-Install-Guide/\
OVS_MultiNode/KeystoneScripts/keystone_endpoints_basic.sh
```

```
chmod +x keystone_basic.sh
chmod +x keystone_endpoints_basic.sh
```

```
./keystone_basic.sh
./keystone_endpoints_basic.sh
```

- Create a simple credential file and load it so you won't be bothered later :

```
nano creds
```

```
#Paste the following:
export OS_TENANT_NAME=admin
export OS_USERNAME=admin
export OS_PASSWORD=admin_pass
export OS_AUTH_URL="http://192.168.100.51:5000/v2.0/"
```

```
# Load it:
source creds
```

- To test Keystone, we use a simple CLI command :
`keystone user-list`

7 Glance

- We Move now to Glance installation :
`apt-get install -y glance`
- Update `/etc/glance/glance-api-paste.ini` with :

```
[filter:authtoken]
paste.filter_factory = keystoneclient.middleware.auth_token:\
    filter_factory
delay_auth_decision = true
auth_host = 10.10.10.51
auth_port = 35357
auth_protocol = http
admin_tenant_name = service
admin_user = glance
admin_password = service_pass
```
- Update the `/etc/glance/glance-registry-paste.ini` with :

```
[filter:authtoken]
paste.filter_factory = keystoneclient.middleware.auth_token:\
    filter_factory
auth_host = 10.10.10.51
auth_port = 35357
auth_protocol = http
admin_tenant_name = service
admin_user = glance
admin_password = service_pass
```
- Update `/etc/glance/glance-api.conf` with :

```
sql_connection = mysql://glanceUser:glancePass@10.10.10.51/glance
```
- And :
`[paste_deploy]`

- ```
flavor = keystone
```
- Update the /etc/glance/glance-registry.conf with :
 

```
sql_connection = mysql://glanceUser:glancePass@10.10.10.51/glance
```
  - And :
 

```
[paste_deploy]
flavor = keystone
```
  - Restart the glance-api and glance-registry services :
 

```
service glance-api restart; service glance-registry restart
```
  - Synchronize the glance database :
 

```
glance-manage db_sync
```
  - To test Glance, upload the cirros cloud image directly from the internet :
 

```
glance image-create --name myFirstImage --is-public true \
--container-format bare --disk-format qcow2 --location \
https://launchpad.net/cirros/trunk/0.3.0/+download/\
cirros-0.3.0-x86_64-disk.img
```
  - Now list the image to see what you have just uploaded :
 

```
glance image-list
```

## 8 Quantum

- Install the Quantum server and the OpenVSwitch package collection :
 

```
apt-get install -y quantum-server
```
- Edit the OVS plugin configuration file /etc/quantum/plugins/openvswitch/ovs\_quantum\_plugin.ini with :
 

```
#Under the database section
[DATABASE]
sql_connection = mysql://quantumUser:quantumPass@10.10.10.51/quantum

#Under the OVS section
[OVS]
tenant_network_type = gre
tunnel_id_ranges = 1:1000
enable_tunneling = True
```
- Edit /etc/quantum/api-paste.ini :
 

```
[filter:authtoken]
paste.filter_factory = keystoneclient.middleware.auth_token:\
 filter_factory
auth_host = 10.10.10.51
auth_port = 35357
auth_protocol = http
```



- ```

admin_tenant_name = service
admin_user = quantum
admin_password = service_pass

```
- Update the /etc/quantum/quantum.conf :

```

[keystone_auth_token]
auth_host = 10.10.10.51
auth_port = 35357
auth_protocol = http
admin_tenant_name = service
admin_user = quantum
admin_password = service_pass
signing_dir = /var/lib/quantum/keystone-signing

```
 - Restart the quantum server :

```

service quantum-server restart

```

9 Nova

- Start by installing nova components :

```

apt-get install -y nova-api nova-cert novnc nova-consoleauth \
nova-scheduler nova-novncproxy nova-doc nova-conductor

```
- Now modify auth_token section in the /etc/nova/api-paste.ini file to this :

```

[filter:auth_token]
paste.filter_factory = keystoneclient.middleware.auth_token:\
    filter_factory
auth_host = 10.10.10.51
auth_port = 35357
auth_protocol = http
admin_tenant_name = service
admin_user = nova
admin_password = service_pass
signing_dirname = /tmp/keystone-signing-nova
# Workaround for https://bugs.launchpad.net/nova/+bug/1154809
auth_version = v2.0

```
- Modify the /etc/nova/nova.conf like this :

```

[DEFAULT]
logdir=/var/log/nova
state_path=/var/lib/nova
lock_path=/run/lock/nova
verbose=True
api_paste_config=/etc/nova/api-paste.ini

```

```
compute_scheduler_driver=nova.scheduler.simple.SimpleScheduler
rabbit_host=10.10.10.51
nova_url=http://10.10.10.51:8774/v1.1/
sql_connection=mysql://novaUser:novaPass@10.10.10.51/nova
root_helper=sudo nova-rootwrap /etc/nova/rootwrap.conf

# Auth
use_deprecated_auth=false
auth_strategy=keystone

# Imaging service
glance_api_servers=10.10.10.51:9292
image_service=nova.image.glance.GlanceImageService

# Vnc configuration
novnc_enabled=true
novncproxy_base_url=http://192.168.100.51:6080/vnc_auto.html
novncproxy_port=6080
vncserver_proxycient_address=10.10.10.51
vncserver_listen=0.0.0.0

# Network settings
network_api_class=nova.network.quantumv2.api.API
quantum_url=http://10.10.10.51:9696
quantum_auth_strategy=keystone
quantum_admin_tenant_name=service
quantum_admin_username=quantum
quantum_admin_password=service_pass
quantum_admin_auth_url=http://10.10.10.51:35357/v2.0
libvirt_vif_driver=nova.virt.libvirt.vif.LibvirtHybridOVSBridgeDriver
linuxnet_interface_driver=nova.network.linux_net.LinuxOVSInterfaceDriver
firewall_driver=nova.virt.libvirt.firewall.IptablesFirewallDriver

#Metadata
service_quantum_metadata_proxy = True
quantum_metadata_proxy_shared_secret = helloOpenStack
metadata_host = 10.10.10.51
metadata_listen = 10.10.10.51
metadata_listen_port = 8775

# Compute #
```

```
compute_driver=libvirt.LibvirtDriver

# Cinder #
volume_api_class=nova.volume.cinder.API
osapi_volume_listen_port=5900
- Synchronize your database :
  nova-manage db sync
- Restart nova-* services :
  cd /etc/init.d/; for i in $( ls nova-* ); do sudo service $i restart; done
- Check for the smiling faces on nova-* services to confirm your installation :
  nova-manage service list
```

10 Cinder

```
- Install the required packages :
  apt-get install -y cinder-api cinder-scheduler cinder-volume \
  iscsitarget open-iscsi iscsitarget-dkms
- Configure the iscsi services :
  sed -i 's/false/true/g' /etc/default/iscsitarget
- Restart the services :
  service iscsitarget start
  service open-iscsi start
- Configure /etc/cinder/api-paste.ini like the following :
  [filter:authtoken]
  paste.filter_factory = keystoneclient.middleware.auth_token:
    filter_factory
  service_protocol = http
  service_host = 192.168.100.51
  service_port = 5000
  auth_host = 10.10.10.51
  auth_port = 35357
  auth_protocol = http
  admin_tenant_name = service
  admin_user = cinder
  admin_password = service_pass
  signing_dir = /var/lib/cinder
- Edit the /etc/cinder/cinder.conf to :
  [DEFAULT]
  rootwrap_config=/etc/cinder/rootwrap.conf
  sql_connection = mysql://cinderUser:cinderPass@10.10.10.51/cinder
```

- ```
api_paste_config = /etc/cinder/api-paste.ini
iscsi_helper=ietadm
volume_name_template = volume-%s
volume_group = cinder-volumes
verbose = True
auth_strategy = keystone
iscsi_ip_address=10.10.10.51
```
- Then, synchronize your database :

```
cinder-manage db sync
```
  - Finally, don't forget to create a volume group and name it cinder-volumes :

```
dd if=/dev/zero of=cinder-volumes bs=1 count=0 seek=2G
losetup /dev/loop2 cinder-volumes
fdisk /dev/loop2
#Type in the followings:
n
p
1
ENTER
ENTER
t
8e
w
```
  - Proceed to create the physical volume then the volume group :

```
pvcreate /dev/loop2
vgcreate cinder-volumes /dev/loop2
```

**Note :** Beware that this volume group gets lost after a system reboot. [Click Here](#) to know how to load it after a reboot

- Restart the cinder services :

```
cd /etc/init.d/; for i in $(ls cinder-*); do sudo service $i restart; done
```
- Verify if cinder services are running :

```
cd /etc/init.d/; for i in $(ls cinder-*); do sudo service $i status; done
```

## 11 Horizon

- To install horizon, proceed like this :

```
apt-get install -y openstack-dashboard memcached
```
- If you don't like the OpenStack ubuntu theme, you can remove the package to disable it :

```
dpkg --purge openstack-dashboard-ubuntu-theme
```

- Reload Apache and memcached :

```
service apache2 restart; service memcached restart
```

# Network Node

## 1 Preparing the Node

- After you install Ubuntu 12.04 or 13.04 Server 64bits, Go in sudo mode :  
`sudo su`
- Add Grizzly repositories [Only for Ubuntu 12.04] :  
`apt-get install -y ubuntu-cloud-keyring`  
`echo deb http://ubuntu-cloud.archive.canonical.com/ubuntu\`  
`precise-updates/grizzly main >> /etc/apt/sources.list.d/grizzly.list`
- Update your system :  
`apt-get update -y`  
`apt-get upgrade -y`  
`apt-get dist-upgrade -y`
- Install ntp service :  
`apt-get install -y ntp`
- Configure the NTP server to follow the controller node :  
`#Comment the ubuntu NTP servers`  
`sed -i 's/server 0.ubuntu.pool.ntp.org/#server 0.ubuntu.pool.ntp.org/g' \`  
`/etc/ntp.conf`  
`sed -i 's/server 1.ubuntu.pool.ntp.org/#server 1.ubuntu.pool.ntp.org/g' \`  
`/etc/ntp.conf`  
`sed -i 's/server 2.ubuntu.pool.ntp.org/#server 2.ubuntu.pool.ntp.org/g' \`  
`/etc/ntp.conf`  
`sed -i 's/server 3.ubuntu.pool.ntp.org/#server 3.ubuntu.pool.ntp.org/g' \`  
`/etc/ntp.conf`  
  
`#Set the network node to follow up your conroller node`  
`sed -i 's/server ntp.ubuntu.com/server 10.10.10.51/g' /etc/ntp.conf`  
  
`service ntp restart`
- Install other services :  
`apt-get install -y vlan bridge-utils`

- Enable IP\_Forwarding :  

```
sed -i 's/#net.ipv4.ip_forward=1/net.ipv4.ip_forward=1/'\n/etc/sysctl.conf
```

  

```
To save you from rebooting, perform the following
sysctl net.ipv4.ip_forward=1
```

## 2 Networking

- 3 NICs must be present :  

```
OpenStack management
auto eth0
iface eth0 inet static
address 10.10.10.52
netmask 255.255.255.0
```

  

```
VM Configuration
auto eth1
iface eth1 inet static
address 10.20.20.52
netmask 255.255.255.0
```

  

```
VM internet Access
auto eth2
iface eth2 inet static
address 192.168.100.52
netmask 255.255.255.0
```

## 3 OpenVSwitch (Part1)

- Install the openVSwitch :  

```
apt-get install -y openvswitch-switch openvswitch-datapath-dkms
```
- Create the bridges :  

```
#br-int will be used for VM integration
ovs-vsctl add-br br-int
```

  

```
#br-ex is used to make to VM accessible from the internet
ovs-vsctl add-br br-ex
```

## 4 Quantum

- Install the Quantum openvswitch agent, l3 agent and dhcp agent :  

```
apt-get -y install quantum-plugin-openvswitch-agent quantum-dhcp-agent \
quantum-l3-agent quantum-metadata-agent
```
- Edit /etc/quantum/api-paste.ini :  

```
[filter:authtoken]
paste.filter_factory = keystoneclient.middleware.auth_token:
 filter_factory
auth_host = 10.10.10.51
auth_port = 35357
auth_protocol = http
admin_tenant_name = service
admin_user = quantum
admin_password = service_pass
```
- Edit the OVS plugin configuration file /etc/quantum/plugins/openvswitch/ovs\_quantum\_plugin.ini with :  

```
#Under the database section
[DATABASE]
sql_connection = mysql://quantumUser:quantumPass@10.10.10.51/quantum

#Under the OVS section
[OVS]
tenant_network_type = gre
tunnel_id_ranges = 1:1000
integration_bridge = br-int
tunnel_bridge = br-tun
local_ip = 10.20.20.52
enable_tunneling = True
```
- Update /etc/quantum/metadata\_agent.ini :  

```
The Quantum user information for accessing the Quantum API.
auth_url = http://10.10.10.51:35357/v2.0
auth_region = RegionOne
admin_tenant_name = service
admin_user = quantum
admin_password = service_pass

IP address used by Nova metadata server
nova_metadata_ip = 10.10.10.51
```



```
TCP Port used by Nova metadata server
nova_metadata_port = 8775

metadata_proxy_shared_secret = helloOpenStack
```

- Make sure that your rabbitMQ IP in /etc/quantum/quantum.conf is set to the controller node :

```
rabbit_host = 10.10.10.51
```

```
#And update the keystone_auth token section

[keystone_auth token]
auth_host = 10.10.10.51
auth_port = 35357
auth_protocol = http
admin_tenant_name = service
admin_user = quantum
admin_password = service_pass
signing_dir = /var/lib/quantum/keystone-signing
```

- Restart all the services :

```
cd /etc/init.d/; for i in $(ls quantum-*);
do sudo service $i restart; done
```

## 5 OpenVSwitch (Part2)

- Edit the eth2 in /etc/network/interfaces to become like this :

```
VM internet Access
auto eth2
iface eth2 inet manual
up ifconfig $IFACE 0.0.0.0 up
up ip link set $IFACE promisc on
down ip link set $IFACE promisc off
down ifconfig $IFACE down
```
- Add the eth2 to the br-ex :

```
#Internet connectivity will be lost after this step but\
this won't affect OpenStack's work
ovs-vsctl add-port br-ex eth2
```

If you want to get internet connection back, you can assign the eth2's IP address to the br-ex in the /etc/network interfaces file.

# Compute Node

## 1 Preparing the Node

- After you install Ubuntu 12.04 or 13.04 Server 64bits, Go in sudo mode :  
`sudo su`
- Add Grizzly repositories [Only for Ubuntu 12.04] :  
`apt-get install -y ubuntu-cloud-keyring`  
`echo deb http://ubuntu-cloud.archive.canonical.com/ubuntu\`  
`precise-updates/grizzly main >> /etc/apt/sources.list.d/grizzly.list`
- Update your system :  
`apt-get update -y`  
`apt-get upgrade -y`  
`apt-get dist-upgrade -y`
- Install ntp service :  
`apt-get install -y ntp`
- Configure the NTP server to follow the controller node :  
`#Comment the ubuntu NTP servers`  
`sed -i 's/server 0.ubuntu.pool.ntp.org/#server 0.ubuntu.pool.ntp.org/g' \`  
`/etc/ntp.conf`  
`sed -i 's/server 1.ubuntu.pool.ntp.org/#server 1.ubuntu.pool.ntp.org/g' \`  
`/etc/ntp.conf`  
`sed -i 's/server 2.ubuntu.pool.ntp.org/#server 2.ubuntu.pool.ntp.org/g' \`  
`/etc/ntp.conf`  
`sed -i 's/server 3.ubuntu.pool.ntp.org/#server 3.ubuntu.pool.ntp.org/g' \`  
`/etc/ntp.conf`  
  
`#Set the compute node to follow up your controller node`  
`sed -i 's/server ntp.ubuntu.com/server 10.10.10.51/g' /etc/ntp.conf`  
  
`service ntp restart`
- Install other services :  
`apt-get install -y vlan bridge-utils`

- Enable IP\_Forwarding :  

```
sed -i 's/#net.ipv4.ip_forward=1/net.ipv4.ip_forward=1/' \
/etc/sysctl.conf
```

  

```
To save you from rebooting, perform the following
sysctl net.ipv4.ip_forward=1
```

## 2 Networking

- Perform the following :  

```
OpenStack management
auto eth0
iface eth0 inet static
address 10.10.10.53
netmask 255.255.255.0
```

  

```
VM Configuration
auto eth1
iface eth1 inet static
address 10.20.20.53
netmask 255.255.255.0
```

## 3 KVM

- make sure that your hardware enables virtualization :  

```
apt-get install -y cpu-checker
kvm-ok
```
- Normally you would get a good response. Now, move to install kvm and configure it :  

```
apt-get install -y kvm libvirt-bin pm-utils
```
- Edit the `cgroup_device_acl` array in the `/etc/libvirt/qemu.conf` file to :  

```
cgroup_device_acl = [
 "/dev/null", "/dev/full", "/dev/zero",
 "/dev/random", "/dev/urandom",
 "/dev/ptmx", "/dev/kvm", "/dev/kgemu",
 "/dev/rtc", "/dev/hpet", "/dev/net/tun"
]
```
- Delete default virtual bridge :  

```
virsh net-destroy default
```

- ```
virsh net-undefine default
```
- Enable live migration by updating `/etc/libvirt/libvirtd.conf` file :

```
listen_tls = 0
listen_tcp = 1
auth_tcp = "none"
```
 - Edit `libvirtd_opts` variable in `/etc/init/libvirt-bin.conf` file :

```
env libvirtd_opts="-d -l"
```
 - Edit `/etc/default/libvirt-bin` file :

```
libvirtd_opts="-d -l"
```
 - Restart the libvirt service to load the new values :

```
service libvirt-bin restart
```

4 OpenVSwitch

- Install the openVSwitch :

```
apt-get install -y openvswitch-switch openvswitch-datapath-dkms
```
- Create the bridges :

```
#br-int will be used for VM integration
ovs-vsctl add-br br-int
```

5 Quantum

- Install the Quantum openvswitch agent :

```
apt-get -y install quantum-plugin-openvswitch-agent
```
- Edit the OVS plugin configuration file `/etc/quantum/plugins/openvswitch/ovs_quantum_plugin.ini` with :

```
#Under the database section
[DATABASE]
sql_connection = mysql://quantumUser:quantumPass@10.10.10.51/quantum

#Under the OVS section
[OVS]
tenant_network_type = gre
tunnel_id_ranges = 1:1000
integration_bridge = br-int
tunnel_bridge = br-tun
local_ip = 10.20.20.53
enable_tunneling = True
```

- Make sure that your rabbitMQ IP in `/etc/quantum/quantum.conf` is set to the controller node :

```
rabbit_host = 10.10.10.51
```

#And update the keystone_auth token section

```
[keystone_auth token]
auth_host = 10.10.10.51
auth_port = 35357
auth_protocol = http
admin_tenant_name = service
admin_user = quantum
admin_password = service_pass
signing_dir = /var/lib/quantum/keystone-signing
```

- Restart all the services :

```
service quantum-plugin-openvswitch-agent restart
```

6 Nova

- Install nova's required components for the compute node :

```
apt-get install -y nova-compute-kvm
```

- Now modify auth token section in the `/etc/nova/api-paste.ini` file to this :

```
[filter:auth token]
paste.filter_factory = keystoneclient.middleware.auth_token:\
    filter_factory
auth_host = 10.10.10.51
auth_port = 35357
auth_protocol = http
admin_tenant_name = service
admin_user = nova
admin_password = service_pass
signing_dirname = /tmp/keystone-signing-nova
# Workaround for https://bugs.launchpad.net/nova/+bug/1154809
auth_version = v2.0
```

- Edit `/etc/nova/nova-compute.conf` file :

```
[DEFAULT]
libvirt_type=kvm
libvirt_ovs_bridge=br-int
libvirt_vif_type=ethernet
libvirt_vif_driver=nova.virt.libvirt.vif.LibvirtHybridOVSBridgeDriver
```

```
libvirt_use_virtio_for_bridges=True
- Modify the /etc/nova/nova.conf like this :
[DEFAULT]
logdir=/var/log/nova
state_path=/var/lib/nova
lock_path=/run/lock/nova
verbose=True
api_paste_config=/etc/nova/api-paste.ini
compute_scheduler_driver=nova.scheduler.simple.SimpleScheduler
rabbit_host=10.10.10.51
nova_url=http://10.10.10.51:8774/v1.1/
sql_connection=mysql://novaUser:novaPass@10.10.10.51/nova
root_helper=sudo nova-rootwrap /etc/nova/rootwrap.conf

# Auth
use_deprecated_auth=false
auth_strategy=keystone

# Imaging service
glance_api_servers=10.10.10.51:9292
image_service=nova.image.glance.GlanceImageService

# Vnc configuration
novnc_enabled=true
novncproxy_base_url=http://192.168.100.51:6080/vnc_auto.html
novncproxy_port=6080
vncserver_proxyclient_address=10.10.10.53
vncserver_listen=0.0.0.0

# Network settings
network_api_class=nova.network.quantumv2.api.API
quantum_url=http://10.10.10.51:9696
quantum_auth_strategy=keystone
quantum_admin_tenant_name=service
quantum_admin_username=quantum
quantum_admin_password=service_pass
quantum_admin_auth_url=http://10.10.10.51:35357/v2.0
libvirt_vif_driver=nova.virt.libvirt.vif.LibvirtHybridOVSBridgeDriver
linuxnet_interface_driver=nova.network.\
    linux_net.LinuxOVSInterfaceDriver
firewall_driver=nova.virt.libvirt.firewall.IptablesFirewallDriver
```

```
#Metadata
service_quantum_metadata_proxy = True
quantum_metadata_proxy_shared_secret = helloOpenStack

# Compute #
compute_driver=libvirt.LibvirtDriver

# Cinder #
volume_api_class=nova.volume.cinder.API
osapi_volume_listen_port=5900
cinder_catalog_info=volume:cinder:internalURL
- Restart nova-* services :
  cd /etc/init.d/; for i in $( ls nova-* );
  do sudo service $i restart; done
- Check for the smiling faces on nova-* services to confirm your installation :
  nova-manage service list
```

Your first VM

To start your first VM, we first need to create a new tenant, user and internal network.

- Create a new tenant :
`keystone tenant-create --name project_one`
- Create a new user and assign the member role to it in the new tenant (keystone role-list to get the appropriate id) :
`keystone user-create --name=user_one --pass=user_one \`
`--tenant-id $put_id_of_project_one --email=user_one@domain.com`
`keystone user-role-add --tenant-id $put_id_of_project_one \`
`--user-id $put_id_of_user_one --role-id $put_id_of_member_role`
- Create a new network for the tenant :
`quantum net-create --tenant-id $put_id_of_project_one net_proj_one`
- Create a new subnet inside the new tenant network :
`quantum subnet-create --tenant-id \`
`$put_id_of_project_one net_proj_one 50.50.1.0/24`
- Create a router for the new tenant :
`quantum router-create --tenant-id \`
`$put_id_of_project_one router_proj_one`
- Add the router to the running l3 agent (if it wasn't automatically added) :
`quantum agent-list (to get the l3 agent ID)`
`quantum l3-agent-router-add $l3_agent_ID router_proj_one`
- Add the router to the subnet :
`quantum router-interface-add $put_router_proj_one_id_here \`
`$put_subnet_id_here`
- Restart all quantum services :
`cd /etc/init.d/; for i in $(ls quantum-*);`
`do sudo service $i restart; done`
- Create an external network with the tenant id belonging to the admin tenant (keystone tenant-list to get the appropriate id) :
`quantum net-create --tenant-id $put_id_of_admin_tenant ext_net \`
`--router:external=True`
- Create a subnet for the floating ips :


```
quantum subnet-create --tenant-id $put_id_of_admin_tenant \
--allocation-pool start=192.168.100.102,end=192.168.100.126 \
--gateway 192.168.100.1 ext_net 192.168.100.100/24 \
--enable_dhcp=False
```

- Set your router's gateway to the external network :

```
quantum router-gateway-set $put_router_proj_one_id_here $put_id_of_ext_net_
```

- Source creds relative to your project one tenant now :

```
nano creds_proj_one
```

```
#Paste the following:
```

```
export OS_TENANT_NAME=project_one
```

```
export OS_USERNAME=user_one
```

```
export OS_PASSWORD=user_one
```

```
export OS_AUTH_URL="http://192.168.100.51:5000/v2.0/"
```

```
source creds_proj_one
```

- Start by allocating a floating ip to the project one tenant :

```
quantum floatingip-create ext_net
```

- Start a VM :

```
nova --no-cache boot --image $id_myFirstImage --flavor 1 my_first_vm
```

- pick the id of the port corresponding to your VM :

```
quantum port-list
```

- Associate the floating IP to your VM :

```
quantum floatingip-associate $put_id_floating_ip $put_id_vm_port
```

That's it! ping your VM and enjoy your OpenStack.

Licensing

OpenStack Grizzly Install Guide is licensed under a Creative Commons Attribution 3.0 Unported License.

To view a copy of this license, visit [http ://creativecommons.org/licenses/by/3.0/deed.en_US](http://creativecommons.org/licenses/by/3.0/deed.en_US).

Credits

This work has been based on :

- Bilel Msekni's Folsom Install guide [https ://github.com/mseknbilel/OpenStack-Folsom-Install-guide](https://github.com/mseknbilel/OpenStack-Folsom-Install-guide)
- OpenStack Grizzly Install Guide (Master Branch) [https ://github.com/mseknbilel/OpenStack-Grizzly-Install-Guide](https://github.com/mseknbilel/OpenStack-Grizzly-Install-Guide)